

✓ Implementation of Reinforce from scratch to play Cartpole-v1

- Environment documentations: https://gymnasium.farama.org/environments/classic_control/cart_pole/

✓ Install and import libraries

```

1 # Import Libraries
2
3 import gymnasium as gym
4 import torch
5 import torch.nn as nn
6 from torch.distributions.categorical import Categorical
7 from torch.distributions.normal import Normal
8 from collections import deque
9 import numpy as np
10
11 # For pushing to hub
12 from huggingface_hub import HfApi, snapshot_download
13 from huggingface_hub.repocard import metadata_eval_result, metadata_save
14 from huggingface_hub import notebook_login
15
16 from pathlib import Path
17 import datetime
18 import json
19 import imageio
20
21 import tempfile
22
23 import os

```

```

1 # Device
2
3 device = 'cuda' if torch.cuda.is_available() else 'cpu'

```

✓ 0. Visualise the observation space and action space

```

1 env = gym.make("CartPole-v1")
2
3 print(f"Number of available actions: {env.action_space.n}")
4 print(f"Sample a random action: {env.action_space.sample()}")
5
6 print(f"Sample a random observation: {env.observation_space.sample()}")

```

```

↗ Number of available actions: 2
Sample a random action: 0
Sample a random observation: [4.745393  2.4669313  0.4147641  0.27528846]

```

✓ 1. Build a policy network using PyTorch

```

1 class Policy(nn.Module):
2     """
3     A policy network.
4
5     Args:
6         s_size (int): The size of 1 state space. \n
7         h_size (int): The number of hidden nodes in the network. \n
8         a_size (int): The number of distinct discrete actions, representing the number of output nodes. \n
9     """
10    def __init__(self, s_size, h_size, a_size):
11        super().__init__()
12
13        self.fc1 = nn.Sequential(nn.Linear(in_features = s_size,
14                                           out_features = h_size),
15                                nn.ReLU())
16
17        self.fc2 = nn.Sequential(nn.Linear(in_features = h_size,
18                                           out_features = a_size))
19
20    def forward(self, x):
21        """
22        The forward propagation of the policy network.

```

```

23
24 Args:
25     x (float tensor): Input to the network representing the observation / state, expected shape: (B, s_size). \n
26
27 Returns:
28     out(float tensor): Output of the network, representing the probability of taking each distinct discrete action, expected shape:
29     """
30     out = self.fc1(x)
31     out = self.fc2(out)
32     out = torch.nn.functional.softmax(out, dim=-1)
33     return out
34
35 def act(self, state):
36     """
37     Sampling of an action.
38
39     Args:
40         state (float tensor): Input to the network representing the observation / state, expected shape: (B, s_size). \n
41
42     Returns:
43         action (int / int tensor): The index of the output nodes of the network, sampled based on the output probability of the network.
44         log_prob (float tensor): The ln of the probability of the action that was sampled based on the output probability of the network.
45     """
46     probs = self.forward(state).cpu()
47     m = Categorical(probs)
48     action = m.sample()
49     if action.shape[0] == 1:
50         return action.item(), m.log_prob(action)
51     else:
52         return action, m.log_prob(action)
53
54
55

```

✓ 2. Implementation of Reinforce algorithm

```

1 def reinforce(policy, optimizer, env, n_episodes, n_steps, device, gamma, print_every):
2
3     """
4     Train a policy network using the reinforce algorithm.
5
6     Args:
7         policy (nn.Module): A policy network. \n
8         optimizer (torch.optim): An optimizer. \n
9         env (gymnasium.env): An environment. \n
10        n_episodes (int): Number of training episodes. \n
11        n_steps (int): Maximum number of steps allowed in an episode. \n
12        device (str): 'cuda' or 'cpu'
13        gamma (float): A discount factor, range from 0 to 1.
14        print_every (int): Number of episode intervals to print the performance of the network. \n
15
16    Returns:
17        scores (list): A list of integer, each element representing the rewards scored in an episode. \n
18
19    """
20
21    # Create variables to store the rewards scored for every episode
22    scores = []
23
24    # This variable store up to rewards scored for every episode up to "print_every" episodes
25    scores_deque = deque(maxlen = print_every)
26
27
28    #####
29    # For each episode #
30    #####
31    for episode in range(1, n_episodes+1):
32
33        # Variable to store values for every step within an episode
34        reward_eps = [] # Reward scored by each step
35        log_prob_eps = [] # ln (prob of the action taken) for each step
36        returns_eps = deque(maxlen = n_steps) # discounted returns scored in each step
37        policy_loss_eps = [] # loss for each step -> ln (prob of the action taken) * discounted return
38
39        # Reset the environment for the beginning of each episode
40        state, info = env.reset()
41
42        #####
43        # For each step #
44        #####

```

```

45     for _ in range(n_steps):
46         # Sample an action using the policy network
47         action, log_prob = policy.act(torch.tensor(state).unsqueeze(0).to(device))
48
49         # Step the environment using the sampled action
50         state, reward, terminated, truncated, info = env.step(action)
51
52         # Store the rewards for this step and the ln (prob of this action)
53         reward_eps.append(reward)
54         log_prob_eps.append(log_prob)
55
56         # Check if this leads to termination or truncation
57         if terminated or truncated:
58             break
59
60     # Sum the rewards scored for the entire episode and store them
61     scores.append(sum(reward_eps))
62     scores_deque.append(sum(reward_eps))
63
64     # Calculate the discounted returns
65     for t in range(len(reward_eps))[:-1]:
66         returns = reward_eps[t] + gamma * returns_eps[0] if len(returns_eps) > 0 else reward_eps[t]
67         returns_eps.appendleft(returns)
68
69     # Normalize the discounted returns
70     eps = np.finfo(np.float32).eps.item() # eps is smallest representatable float
71     returns_eps = torch.tensor(returns_eps) # Convert into torch tensor for later calculation
72     returns_eps = (returns_eps - returns_eps.mean()) / (returns_eps.std() + eps)
73
74     # Calculate the loss
75     for log_prob, returns in zip(log_prob_eps, returns_eps):
76         policy_loss_eps.append(-log_prob * returns) # torch tensor * torch tensor
77     loss = torch.cat(policy_loss_eps).sum() # cat -> makes into one torch tensor, sum() -> summation
78
79     # Backward propagation and gradient descent
80     optimizer.zero_grad()
81     loss.backward()
82     optimizer.step()
83
84     # Print information
85     if episode % print_every == 0:
86         print(f"Current Episode: {episode} | Average reward: {np.mean(scores_deque)}")
87
88     return scores
89
90
91
92
93
94
95

```

✓ 3. Train the policy network using reinforce algorithm

```

1 #0. Create device
2 device = 'cuda' if torch.cuda.is_available() else 'cpu'
3
4 # 1. Create hyperparameters
5 cartpole_hyperparameters = {
6     "h_size": 16,
7     "n_training_episodes": 1000,
8     "n_evaluation_episodes": 10,
9     "max_t": 1000,
10    "gamma": 1.0,
11    "lr": 0.01,
12    "env_id": 'CartPole-v1',
13    "state_space": 4,
14    "action_space": 2,
15 }
16
17 # 2. Create environment
18 env = gym.make(cartpole_hyperparameters['env_id'])
19
20 # 3. Create policy network
21 cartpole_policy = Policy(cartpole_hyperparameters['state_space'],
22                          cartpole_hyperparameters['h_size'],
23                          cartpole_hyperparameters['action_space']).to(device)
24
25 # 4. Create optimizer
26 optimizer = torch.optim.Adam(cartpole_policy.parameters()),

```

```

27         lr = cartpole_hyperparameters['lr'])
28
29 # 5. Training loop
30 scores = reinforce(policy = cartpole_policy,
31                    optimizer = optimizer,
32                    env = env,
33                    n_episodes = cartpole_hyperparameters['n_training_episodes'],
34                    n_steps = cartpole_hyperparameters['max_t'],
35                    device = device,
36                    gamma = cartpole_hyperparameters['gamma'],
37                    print_every = 100)
38

```

```

↻ Current Episode: 100 | Average reward: 44.33
Current Episode: 200 | Average reward: 308.61
Current Episode: 300 | Average reward: 322.93
Current Episode: 400 | Average reward: 388.92
Current Episode: 500 | Average reward: 352.5
Current Episode: 600 | Average reward: 459.3
Current Episode: 700 | Average reward: 476.28
Current Episode: 800 | Average reward: 500.0
Current Episode: 900 | Average reward: 461.11
Current Episode: 1000 | Average reward: 379.94

```

✓ 4. Evaluate the agent

```

1 def evaluate_agent(n_eval_episodes, n_steps, policy, env, device):
2
3     """
4     Evaluate the performance of an agent by calculating the mean and standard deviation rewards over n_eval_episodes of episodes.
5
6     Args:
7         n_eval_episodes (int): Number of evaluation episodes. \n
8         n_steps (int): Maximum number of steps allowed in an episode. \n
9         policy (nn.Module): A policy network. \n
10        env (gymnasium.env): An environment. \n
11        device (str): 'cuda' or 'cpu'. \n
12
13    Returns:
14        mean_reward (float): Mean reward scored across the evaluated episodes.
15        std_reward (float): Standard deviation reward scored across the evaluated episodes.
16    """
17
18
19    rewards_across_episodes = [] # Contains rewards scored in each episode
20
21    #####
22    # For each episode #
23    #####
24    for episode in range(n_eval_episodes):
25
26        # To store reward scored in each step
27        rewards = []
28
29        # Reset the environment
30        state, info = env.reset()
31
32        #####
33        # For each step #
34        #####
35        for step in range(n_steps):
36
37            # Sample an action
38            action, _ = policy.act(torch.tensor(state).unsqueeze(0).to(device))
39
40            # Step the environment by taking the action
41            state, reward, terminated, truncated, info = env.step(action)
42
43            # Store the reward scored in this step
44            rewards.append(reward)
45
46            # Check if truncated or terminated
47            if truncated or terminated:
48                break
49
50            # Sum the reward scored in the entire episode and store it
51            rewards_across_episodes.append(sum(rewards))
52
53    # Calculate the mean and standard deviation
54    mean_reward = np.array(rewards_across_episodes).mean()
55    std_reward = np.array(rewards_across_episodes).std()

```

```
56 return mean_reward, std_reward
57
```

```
1 mean_reward, std_reward = evaluate_agent(n_eval_episodes = cartpole_hyperparameters['n_evaluation_episodes'],
2                                         n_steps = cartpole_hyperparameters['max_t'],
3                                         policy = cartpole_policy,
4                                         env = env,
5                                         device = device)
6
7 print(f"Mean reward: {mean_reward}, standard deviation: {std_reward}")
```

↗ Mean reward: 500.0, standard deviation: 0.0

5. Push to Hub

- Code source: https://colab.research.google.com/github/wengti/Reinforcement-Learning-Tutorial/blob/main/notebooks/unit4/unit4.ipynb#scrollTo=LIVsvlW_8tcw
- Create a write token here: <https://huggingface.co/settings/tokens/new?tokenType=write>

```
1 def record_video(env, policy, out_directory, device, fps=30):
2     """
3     Generate a replay video of the agent
4     :param env
5     :param Qtable: Qtable of our agent
6     :param out_directory
7     :param fps: how many frame per seconds (with taxi-v3 and frozenlake-v1 we use 1)
8     """
9     images = []
10    state, info = env.reset()
11    terminated = False
12    truncated = False
13    img = env.render()
14    images.append(img)
15    while not terminated and not truncated:
16        # Take the action (index) that have the maximum expected future reward given that state
17        action, _ = policy.act(torch.tensor(state).unsqueeze(0).to(device))
18        state, reward, terminated, truncated, info = env.step(action) # We directly put next_state = state for recording logic
19        img = env.render()
20        images.append(img)
21    imageio.mimsave(out_directory, [np.array(img) for i, img in enumerate(images)], fps=fps)
```

```
1 def push_to_hub(repo_id,
2                 model,
3                 hyperparameters,
4                 eval_env,
5                 video_fps=30
6                 ):
7     """
8     Evaluate, Generate a video and Upload a model to Hugging Face Hub.
9     This method does the complete pipeline:
10    - It evaluates the model
11    - It generates the model card
12    - It generates a replay video of the agent
13    - It pushes everything to the Hub
14
15    :param repo_id: repo_id: id of the model repository from the Hugging Face Hub
16    :param model: the pytorch model we want to save
17    :param hyperparameters: training hyperparameters
18    :param eval_env: evaluation environment
19    :param video_fps: how many frame per seconds to record our video replay
20    """
21
22    _, repo_name = repo_id.split("/")
23    api = HfApi()
24
25    # Step 1: Create the repo
26    repo_url = api.create_repo(
27        repo_id=repo_id,
28        exist_ok=True,
29    )
30
31    with tempfile.TemporaryDirectory() as tmpdirname:
32        local_directory = Path(tmpdirname)
33
34        # Step 2: Save the model
35        torch.save(model, local_directory / "model.pt")
36
```

```

37 # Step 3: Save the hyperparameters to JSON
38 with open(local_directory / "hyperparameters.json", "w") as outfile:
39     json.dump(hyperparameters, outfile)
40
41 # Step 4: Evaluate the model and build JSON
42 mean_reward, std_reward = evaluate_agent(hyperparameters["n_evaluation_episodes"],
43                                         hyperparameters["max_t"],
44                                         model,
45                                         eval_env,
46                                         'cuda')
47
48 # Get datetime
49 eval_datetime = datetime.datetime.now()
50 eval_form_datetime = eval_datetime.isoformat()
51
52 evaluate_data = {
53     "env_id": hyperparameters["env_id"],
54     "mean_reward": mean_reward,
55     "n_evaluation_episodes": hyperparameters["n_evaluation_episodes"],
56     "eval_datetime": eval_form_datetime,
57 }
58
59 # Write a JSON file
60 with open(local_directory / "results.json", "w") as outfile:
61     json.dump(evaluate_data, outfile)
62
63 # Step 5: Create the model card
64 env_name = hyperparameters["env_id"]
65 env_id = env_name
66
67 metadata = {}
68 metadata["tags"] = [
69     env_name,
70     "reinforce",
71     "reinforcement-learning",
72     "custom-implementation",
73     "deep-rl-class"
74 ]
75
76 # Add metrics
77 eval = metadata_eval_result(
78     model_pretty_name=repo_name,
79     task_pretty_name="reinforcement-learning",
80     task_id="reinforcement-learning",
81     metrics_pretty_name="mean_reward",
82     metrics_id="mean_reward",
83     metrics_value=f"{mean_reward:.2f} +/- {std_reward:.2f}",
84     dataset_pretty_name=env_name,
85     dataset_id=env_name,
86 )
87
88 # Merges both dictionaries
89 metadata = {**metadata, **eval}
90
91 model_card = f"""
92 # **Reinforce** Agent playing **{env_id}**
93 This is a trained model of a **Reinforce** agent playing **{env_id}** .
94 To learn to use this model and train yours check Unit 4 of the Deep Reinforcement Learning Course: https://huggingface.co/deep-rl
95 """
96
97 readme_path = local_directory / "README.md"
98 readme = ""
99 if readme_path.exists():
100     with readme_path.open("r", encoding="utf8") as f:
101         readme = f.read()
102 else:
103     readme = model_card
104
105 with readme_path.open("w", encoding="utf-8") as f:
106     f.write(readme)
107
108 # Save our metrics to Readme metadata
109 metadata_save(readme_path, metadata)
110
111 # Step 6: Record a video
112 video_path = local_directory / "replay.mp4"
113 record_video(eval_env, model, video_path, 'cuda', video_fps)
114
115 # Step 7. Push everything to the Hub
116 api.upload_folder(
117     repo_id=repo_id,
118     folder_path=local_directory,
119     path_in_repo=".",

```

```

119     )
120
121     print(f"Your model is pushed to the Hub. You can view your model here: {repo_url}")

```

```

1 # Login to hugging face with a write token
2
3 notebook_login()

```



```

1 # Create a new environment with render mode (needed for video recording)
2 eval_env = gym.make(cartpole_hyperparameters['env_id'], render_mode = 'rgb_array')
3
4 # Push to Hub
5 push_to_hub(repo_id = "wengti0608/Reinforce-Cartpole-v1-attempt2",
6             model = cartpole_policy,
7             hyperparameters = cartpole_hyperparameters,
8             eval_env = eval_env)

```

WARNING:imageio_ffmpeg:IMAGEIO FFMPEG_WRITER WARNING: input image is not divisible by macro_block_size=16, resizing from (600, 400)
 Uploading...: 100% 3.78k/3.78k [00:01<00:00, 18.8kB/s]
 Your model is pushed to the Hub. You can view your model here: <https://huggingface.co/wengti0608/Reinforce-Cartpole-v1-attempt2>

✓ Practice: Application of Reinforce to play Continuous Mountain Car

- Environment documentation: https://gymnasium.farama.org/environments/classic_control/mountain_car/
- Continuous Mountain Car agent takes continuous actions. Therefore, the policy network is slightly modified to output mean and standard deviation instead of index of discrete actions.

✓ 0. Visualize Environment

```

1 # Visualize the environment
2
3 env = gym.make("MountainCarContinuous-v0")
4
5 print(f"Randomly sample an action: {env.action_space.sample()}")
6 print(f"Randomly sample a state: {env.observation_space.sample()}")
7

```

Randomly sample an action: [0.9539736]
 Randomly sample a state: [-0.00267963 -0.00615319]

✓ 1. Create the policy network

```

1 class Policy(nn.Module):
2     """
3     A policy network.
4
5     Args:
6         s_size (int): The size of 1 state space. \n
7         h_size (int): The number of hidden nodes in the network. \n
8         a_size (int): Number of continuous actions. \n
9     """
10    def __init__(self, s_size, h_size, a_size):
11        super().__init__()
12
13        self.fc1 = nn.Sequential(nn.Linear(in_features = s_size,
14                                           out_features = h_size),
15                                nn.ReLU())
16
17        self.fc2 = nn.Sequential(nn.Linear(in_features = h_size,
18                                           out_features = h_size),
19                                nn.ReLU())
20
21        self.mean_head = nn.Linear(in_features = h_size,
22                                   out_features = a_size)
23
24        self.log_std_head = nn.Linear(in_features = h_size,
25                                      out_features = a_size)
26
27    def forward(self, x):

```

```

28     """
29     The forward propagation of the policy network.
30
31     Args:
32         x (float tensor): Input to the network representing the observation / state, expected shape: (B, s_size). \n
33
34     Returns:
35         mean (float tensor): Mean of a Normal Distribution, (B, a_size). \n
36         std (float tensor): Standard deviation of a Normal Distribution, (B, a_size). \n
37     """
38     out = self.fc1(x)
39     out = self.fc2(out)
40
41     mean = self.mean_head(out)
42     log_std = self.log_std_head(out)
43     std = torch.exp(log_std)
44
45     return mean, std
46
47 def act(self, state):
48     """
49     Sampling of an action.
50
51     Args:
52         state (float tensor): Input to the network representing the observation / state, expected shape: (B, s_size). \n
53
54     Returns:
55         action_clipped (float): The value of the action taken, in the range of -1 to 1 \n
56         log_prob (float tensor): The ln of the probability of the action that was sampled based on the output probability of the network. \n
57
58     """
59     mean, std = self.forward(state)
60     mean = mean.cpu() # torch.tensor, (1,1)
61     std = std.cpu() # torch.tensor, (1,1)
62
63     m = Normal(mean, std)
64
65     action = m.sample() # torch.tensor, (1,1)
66     action_clipped = torch.clamp(action, -1, 1).item() #float
67
68     log_prob = m.log_prob(action)[0] # torch.tensor, (1,)
69
70     return action_clipped, log_prob
71
72

```

✓ 2. Implement reinforce algorithm from scratch

```

1 def reinforce(policy, optimizer, env, n_episodes, n_steps, device, gamma, print_every):
2
3     """
4     Train a policy network using the reinforce algorithm.
5
6     Args:
7         policy (nn.Module): A policy network. \n
8         optimizer (torch.optim): An optimizer. \n
9         env (gymnasium.env): An environment. \n
10        n_episodes (int): Number of training episodes. \n
11        n_steps (int): Maximum number of steps allowed in an episode. \n
12        device (str): 'cuda' or 'cpu'
13        gamma (float): A discount factor, range from 0 to 1.
14        print_every (int): Number of episode intervals to print the performance of the network. \n
15
16    Returns:
17        scores (list): A list of integer, each element representing the rewards scored in an episode. \n
18
19    """
20
21    # Create variables to store the rewards scored for every episode
22    scores = []
23
24    # This variable store up to rewards scored for every episode up to "print_every" episodes
25    scores_deque = deque(maxlen = print_every)
26    len_deque = deque(maxlen = print_every)
27
28
29    #####
30    # For each episode #
31    #####
32    for episode in range(1, n_episodes+1):

```



```

33
34 # Variable to store values for every step within an episode
35 reward_eps = [] # Reward scored by each step
36 log_prob_eps = [] # ln (prob of the action taken) for each step
37 returns_eps = deque(maxlen = n_steps) # discounted returns scored in each step
38 policy_loss_eps = [] # loss for each step -> ln (prob of the action taken) * discounted return
39
40 # Reset the environment for the beginning of each episode
41 state, info = env.reset()
42
43 #####
44 # For each step #
45 #####
46 for _ in range(n_steps):
47
48     # Sample an action using the policy network
49     action, log_prob = policy.act(torch.tensor(state).unsqueeze(0).to(device))
50
51     # Step the environment using the sampled action
52     state, reward, terminated, truncated, info = env.step(np.array([action]))
53
54     # Store the rewards for this step and the ln (prob of this action)
55     reward_eps.append(reward)
56     log_prob_eps.append(log_prob)
57
58     # Check if this leads to termination or truncation
59     if terminated or truncated:
60         break
61
62     # Sum the rewards scored for the entire episode and store them
63     scores.append(sum(reward_eps))
64     scores_deque.append(sum(reward_eps))
65     len_deque.append(len(reward_eps))
66
67     # Calculate the discounted returns
68     for t in range(len(reward_eps))[:-1]:
69         returns = reward_eps[t] + gamma * returns_eps[0] if len(returns_eps) > 0 else reward_eps[t]
70         returns_eps.appendleft(returns)
71
72     # Normalize the discounted returns
73     eps = np.finfo(np.float32).eps.item() # eps is smallest representatable float
74     returns_eps = torch.tensor(returns_eps) # Convert into torch tensor for later calculation
75     returns_eps = (returns_eps - returns_eps.mean()) / (returns_eps.std() + eps)
76
77     # Calculate the loss
78     for log_prob, returns in zip(log_prob_eps, returns_eps):
79         policy_loss_eps.append(-log_prob * returns) # torch tensor * torch tensor
80     loss = torch.cat(policy_loss_eps).sum() # cat -> makes into one torch tensor, sum() -> summation
81
82     # Backward propagation and gradient descent
83     optimizer.zero_grad()
84     loss.backward()
85     optimizer.step()
86
87     # Print information
88     if episode % print_every == 0:
89         print(f"Current Episode: {episode} | Average reward: {np.mean(scores_deque)} | Average length of ep: {np.mean(len_deque)}")
90
91 return scores
92
93
94
95
96
97

```

✓ 3. Training

```

1 # 0. Device
2 device = 'cuda' if torch.cuda.is_available() else 'cpu'
3
4 # 1. Hyperparameters
5 car_hyperparameters = {
6     "h_size": 64,
7     "n_training_episodes": 2000,
8     "n_evaluation_episodes": 10,
9     "max_t": 999,
10    "gamma": 0.9,
11    "lr": 0.00001,
12    "env_id": 'MountainCarContinuous-v0',

```

```

13     "state_space": 2,
14     "action_space": 1,
15 }
16
17 # 2. Build a Policy Network
18 car_policy = Policy(car_hyperparameters['state_space'],
19                     car_hyperparameters['h_size'],
20                     car_hyperparameters['action_space']).to(device)
21
22 # 3. Create environment
23 env = gym.make(car_hyperparameters['env_id'])
24
25 # 4. Train the network
26 optimizer = torch.optim.Adam(car_policy.parameters(),
27                               lr = car_hyperparameters['lr'])
28
29 scores = reinforce(policy = car_policy,
30                   optimizer = optimizer,
31                   env = env,
32                   n_episodes = car_hyperparameters['n_training_episodes'],
33                   n_steps = car_hyperparameters['max_t'],
34                   device = device,
35                   gamma = car_hyperparameters['gamma'],
36                   print_every = 100)
37

```

 Show hidden output

4. Evaluation

```

1 def evaluate_agent(n_eval_episodes, n_steps, policy, env, device):
2
3     """
4     Evaluate the performance of an agent by calculating the mean and standard deviation rewards over n_eval_episodes of episodes.
5
6     Args:
7         n_eval_episodes (int): Number of evaluation episodes. \n
8         n_steps (int): Maximum number of steps allowed in an episode. \n
9         policy (nn.Module): A policy network. \n
10        env (gymnasium.env): An environment. \n
11        device (str): 'cuda' or 'cpu'. \n
12
13    Returns:
14        mean_reward (float): Mean reward scored across the evaluated episodes.
15        std_reward (float): Standard deviation reward scored across the evaluated episodes.
16    """
17
18
19    rewards_across_episodes = [] # Contains rewards scored in each episode
20
21    #####
22    # For each episode #
23    #####
24    for episode in range(n_eval_episodes):
25
26        # To store reward scored in each step
27        rewards = []
28
29        # Reset the environment
30        state, info = env.reset()
31
32        #####
33        # For each step #
34        #####
35        for step in range(n_steps):
36
37            # Sample an action
38            action, _ = policy.act(torch.tensor(state).unsqueeze(0).to(device))
39
40            # Step the environment by taking the action
41            state, reward, terminated, truncated, info = env.step(np.array([action]))
42
43            # Store the reward scored in this step
44            rewards.append(reward)
45
46            # Check if truncated or terminated
47            if truncated or terminated:
48                break
49
50        # Sum the reward scored in the entire episode and store it

```

```

51     rewards_across_episodes.append(sum(rewards))
52
53 # Calculate the mean and standard deviation
54 mean_reward = np.array(rewards_across_episodes).mean()
55 std_reward = np.array(rewards_across_episodes).std()
56 return mean_reward, std_reward
57

```

```

1 mean_reward, std_reward = evaluate_agent(n_eval_episodes = car_hyperparameters['n_evaluation_episodes'],
2                                         n_steps = car_hyperparameters['max_t'],
3                                         policy = car_policy,
4                                         env = env,
5                                         device = device)
6
7 print(f"Mean reward: {mean_reward}, Std reward: {std_reward}")

```

➡ Mean reward: -46.63831101175287, Std reward: 1.2278728245384134

5. Push to Hub

```

1 def record_video(env, policy, out_directory, device, fps=30):
2     """
3     Generate a replay video of the agent
4     :param env
5     :param Qtable: Qtable of our agent
6     :param out_directory
7     :param fps: how many frame per seconds (with taxi-v3 and frozenlake-v1 we use 1)
8     """
9     images = []
10    state, info = env.reset()
11    terminated = False
12    truncated = False
13    img = env.render()
14    images.append(img)
15    while not terminated and not truncated:
16        # Take the action (index) that have the maximum expected future reward given that state
17        action, _ = policy.act(torch.tensor(state).unsqueeze(0).to(device))
18        state, reward, terminated, truncated, info = env.step(np.array([action])) # We directly put next_state = state for recording log
19        img = env.render()
20        images.append(img)
21    imageio.mimsave(out_directory, [np.array(img) for i, img in enumerate(images)], fps=fps)

```

```

1 def push_to_hub(repo_id,
2                 model,
3                 hyperparameters,
4                 eval_env,
5                 video_fps=30
6                 ):
7     """
8     Evaluate, Generate a video and Upload a model to Hugging Face Hub.
9     This method does the complete pipeline:
10    - It evaluates the model
11    - It generates the model card
12    - It generates a replay video of the agent
13    - It pushes everything to the Hub
14
15    :param repo_id: repo_id: id of the model repository from the Hugging Face Hub
16    :param model: the pytorch model we want to save
17    :param hyperparameters: training hyperparameters
18    :param eval_env: evaluation environment
19    :param video_fps: how many frame per seconds to record our video replay
20    """
21
22    _, repo_name = repo_id.split("/")
23    api = HfApi()
24
25    # Step 1: Create the repo
26    repo_url = api.create_repo(
27        repo_id=repo_id,
28        exist_ok=True,
29    )
30
31    with tempfile.TemporaryDirectory() as tmpdirname:
32        local_directory = Path(tmpdirname)
33
34        # Step 2: Save the model
35        torch.save(model, local_directory / "model.pt")
36

```

```

37 # Step 3: Save the hyperparameters to JSON
38 with open(local_directory / "hyperparameters.json", "w") as outfile:
39     json.dump(hyperparameters, outfile)
40
41 # Step 4: Evaluate the model and build JSON
42 mean_reward, std_reward = evaluate_agent(hyperparameters["n_evaluation_episodes"],
43                                         hyperparameters["max_t"],
44                                         model,
45                                         eval_env,
46                                         'cuda')
47
48 # Get datetime
49 eval_datetime = datetime.datetime.now()
50 eval_form_datetime = eval_datetime.isoformat()
51
52 evaluate_data = {
53     "env_id": hyperparameters["env_id"],
54     "mean_reward": mean_reward,
55     "n_evaluation_episodes": hyperparameters["n_evaluation_episodes"],
56     "eval_datetime": eval_form_datetime,
57 }
58
59 # Write a JSON file
60 with open(local_directory / "results.json", "w") as outfile:
61     json.dump(evaluate_data, outfile)
62
63 # Step 5: Create the model card
64 env_name = hyperparameters["env_id"]
65 env_id = env_name
66
67 metadata = {}
68 metadata["tags"] = [
69     env_name,
70     "reinforce",
71     "reinforcement-learning",
72     "custom-implementation",
73     "deep-rl-class"
74 ]
75
76 # Add metrics
77 eval = metadata_eval_result(
78     model_pretty_name=repo_name,
79     task_pretty_name="reinforcement-learning",
80     task_id="reinforcement-learning",
81     metrics_pretty_name="mean_reward",
82     metrics_id="mean_reward",
83     metrics_value=f"{mean_reward:.2f} +/- {std_reward:.2f}",
84     dataset_pretty_name=env_name,
85     dataset_id=env_name,
86 )
87
88 # Merges both dictionaries
89 metadata = {**metadata, **eval}
90
91 model_card = f"""
92 # **Reinforce** Agent playing **{env_id}**
93 This is a trained model of a **Reinforce** agent playing **{env_id}** .
94 To learn to use this model and train yours check Unit 4 of the Deep Reinforcement Learning Course: https://huggingface.co/deep-rl
95 """
96
97 readme_path = local_directory / "README.md"
98 readme = ""
99 if readme_path.exists():
100     with readme_path.open("r", encoding="utf8") as f:
101         readme = f.read()
102 else:
103     readme = model_card
104
105 with readme_path.open("w", encoding="utf-8") as f:
106     f.write(readme)
107
108 # Save our metrics to Readme metadata
109 metadata_save(readme_path, metadata)
110
111 # Step 6: Record a video
112 video_path = local_directory / "replay.mp4"
113 record_video(eval_env, model, video_path, 'cuda', video_fps)
114
115 # Step 7. Push everything to the Hub
116 api.upload_folder(
117     repo_id=repo_id,
118     folder_path=local_directory,
119     path_in_repo=".",

```

```

119     )
120
121     print(f"Your model is pushed to the Hub. You can view your model here: {repo_url}")

```

```

1 # Login to HuggingFace Hub
2 notebook_login()

```

```

1 # Create an evaluation environment
2 eval_env = gym.make(car_hyperparameters['env_id'], render_mode = 'rgb_array')
3
4 # Push to Hub
5 push_to_hub(repo_id = "wengti0608/Reinforce-MountainCarContinuous-v0-attempt1",
6             model = car_policy,
7             hyperparameters = car_hyperparameters,
8             eval_env = eval_env)

```

⚠ WARNING:imageio_ffmpeg:IMAGEIO_FFMPEG_WRITER WARNING: input image is not divisible by macro_block_size=16, resizing from (600, 400)
 Uploading...: 100% 127k/127k [00:01<00:00, 637kB/s]
 Your model is pushed to the Hub. You can view your model here: <https://huggingface.co/wengti0608/Reinforce-MountainCarContinuous-v0>

✓ Solving Mountain Car with SAC

- Reinforce did not manage to solve Continuous Mountain Car. Therefore, SAC is used instead.

```
1 !pip install stable-baselines3==2.0.0a5
```

⚡ Show hidden output

```
1 !pip install huggingface_sb3
```

⚡ Show hidden output

✓ 1. Training

```

1 import gymnasium as gym
2 from stable_baselines3.sac import SAC
3 from stable_baselines3.common.callbacks import EvalCallback
4 from stable_baselines3.common.monitor import Monitor
5 from stable_baselines3.common.evaluation import evaluate_policy
6
7
8 # 1. Make environment
9 env = gym.make("MountainCarContinuous-v0")
10
11 # 2. Create a SAC model
12
13 # The hyperparameter is provided here:
14 # https://huggingface.co/sb3/sac-MountainCarContinuous-v0
15 policy_kwargs = {'log_std_init': -3.67,
16                 'net_arch': [64, 64]}
17
18 model = SAC(batch_size = 512,
19            buffer_size = 50000,
20            ent_coef = 0.1,
21            gamma = 0.9999,
22            gradient_steps = 32,
23            learning_rate = 0.0003,
24            learning_starts = 0,
25            policy = 'MlpPolicy',
26            policy_kwargs = policy_kwargs,
27            tau = 0.01,
28            train_freq = 32,
29            use_sde = True,
30            env = env)
31
32 # 3. Train the model
33
34 # As SAC does not output rollout on its own
35 # A callback is manually created...
36 eval_env = Monitor(gym.make("MountainCarContinuous-v0", render_mode = "rgb_array"))
37
38 eval_callback = EvalCallback(
39     eval_env,

```

```

40     best_model_save_path = './logs/SAC',
41     log_path = './logs/SAC',
42     eval_freq = 1e3,
43     deterministic = True,
44     render = False
45 )
46
47 # Training
48 model.learn(total_timesteps = 5e4, callback = eval_callback)
49
50 # 4. Save the model
51 model_name = "SAC-MountainCarContinuous-v0"
52 model.save(model_name)
53

```

 [Show hidden output](#)

2. Evaluation

```

1 #1. Create an evaluation environment
2 eval_env = Monitor(gym.make("MountainCarContinuous-v0", render_mode = "rgb_array"))
3
4 #2. Evaluate policy
5 mean_reward, std_reward = evaluate_policy(model = model,
6                                           env = eval_env,
7                                           n_eval_episodes = 10,
8                                           deterministic = True)
9
10 print(f"The mean reward: {mean_reward} | The standard deviation: {std_reward}")

```

 The mean reward: 94.67033819999999 | The standard deviation: 0.2572431881083725

3. Push to Hub

```

1 from huggingface_hub import notebook_login
2
3 notebook_login()

```



```

1 from stable_baselines3.common.vec_env import DummyVecEnv
2 from huggingface_sb3 import package_to_hub
3
4 package_to_hub(model = model,
5               model_name = model_name,
6               model_architecture = "SAC",
7               env_id = "MountainCarContinuous-v0",
8               eval_env = DummyVecEnv([lambda: Monitor(gym.make("MountainCarContinuous-v0", render_mode = "rgb_array"))]),
9               repo_id = "wengti0608/SAC-MountainCarContinuous-v0",
10              commit_message = "First Commit")

```

 [Show hidden output](#)